

Regression Using a Neural Network

In this Assignment #4, we will implement a feedforward neural network from scratch to predict cement strength using the given dataset.

Names and IDs:

Abdelrahman Attia Abdelrahman (20206128)

Shrouk Ashraf Ramdan (20206131)

Ahmed Nasr Hassan (20206129)

First: Let's load and preprocess our data

We start by loading the data from the provided excel file, separating the features and target, and normalizing the features.

```
import pandas as pd
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import MinMaxScaler
import numpy as np

# Loading our dataset
file_path = 'concrete_data.xlsx'
data = pd.read_excel(file_path)

# Separating features and target
features = data.iloc[:, :-1].values # Getting the first 4 columns
target = data.iloc[:, -1].values    # Getting the last column

# Normalizing features
scaler = MinMaxScaler()
normalized_features = scaler.fit_transform(features)

# Splitting into training and testing sets
X_train, X_test, y_train, y_test =
train_test_split(normalized_features, target, test_size=0.25,
random_state=42)

# Checking shapes of train and test
print(f"Training features shape: {X_train.shape}")
print(f"Testing features shape: {X_test.shape}")

Training features shape: (525, 4)
Testing features shape: (175, 4)
```

Second: Let's define our Neural Network

We will create a custom `NeuralNetwork` class that handles initialization, forward propagation, backward propagation, prediction, and error calculation. This way is better than standalone functions.

```
class NeuralNetwork:

    def __init__(self, num_features, num_neurons_in_hidden_layer,
num_predictions, learning_rate):
        # Initialize the number of features, neurons in the hidden
layer, and output predictions
        self.num_features = num_features
        self.num_neurons_in_hidden_layer = num_neurons_in_hidden_layer
        self.num_predictions = num_predictions
        self.learning_rate = learning_rate

        ## We can use either ways for initialization ##

        # Random initialization of weights and biases
        self.weights_from_input_to_hidden =
np.random.rand(self.num_features, self.num_neurons_in_hidden_layer)
        self.bias_for_hidden_layer =
np.random.rand(self.num_neurons_in_hidden_layer)
        self.weights_from_hidden_to_output =
np.random.rand(self.num_neurons_in_hidden_layer, self.num_predictions)
        self.bias_for_output_layer =
np.random.rand(self.num_predictions)

        # Xavier initialization for weights
        #hidden_layer_limit = np.sqrt(6 / (self.num_features +
self.num_neurons_in_hidden_layer))
        #self.weights_from_input_to_hidden = np.random.uniform(-
hidden_layer_limit, hidden_layer_limit, (self.num_features,
self.num_neurons_in_hidden_layer))
        #self.bias_for_hidden_layer =
np.zeros(self.num_neurons_in_hidden_layer)

        #output_layer_limit = np.sqrt(6 /
(self.num_neurons_in_hidden_layer + self.num_predictions))
        #self.weights_from_hidden_to_output = np.random.uniform(-
output_layer_limit, output_layer_limit,
(self.num_neurons_in_hidden_layer, self.num_predictions))
        #self.bias_for_output_layer = np.zeros(self.num_predictions)

    def sigmoid(self, x):
```

```

        return 1 / (1 + np.exp(-x))

def sigmoid_derivative(self, x):
    return x * (1 - x)

def forward_prop(self, inputs):
    # Input to hidden layer
    self.hidden_layer_input = np.dot(inputs,
self.weights_from_input_to_hidden) + self.bias_for_hidden_layer
    self.hidden_layer_output =
self.sigmoid(self.hidden_layer_input)

    # Hidden to output layer
    self.output_layer_input = np.dot(self.hidden_layer_output,
self.weights_from_hidden_to_output) + self.bias_for_output_layer

    # Store for use in backward propagation
    self.output = self.output_layer_input # Linear output for
regression

    return self.output

def backward_prop(self, inputs, actual_output, predicted_output):

    # Calculate output error
    error = actual_output - predicted_output

    # Compute gradients for output layer
    output_layer_gradient = -2 * error # Derivative of loss with
respect to the output
    hidden_to_output_weight_update =
np.outer(self.hidden_layer_output, output_layer_gradient)
    output_bias_update = output_layer_gradient

    # Compute gradients for hidden layer
    hidden_layer_error = output_layer_gradient @
self.weights_from_hidden_to_output.T
    hidden_layer_gradient = hidden_layer_error *
self.sigmoid_derivative(self.hidden_layer_output)
    input_to_hidden_weight_update = np.outer(inputs,
hidden_layer_gradient)
    hidden_bias_update = hidden_layer_gradient

    # Update weights and biases
    self.weights_from_hidden_to_output -= self.learning_rate *

```

```

hidden_to_output_weight_update
    self.bias_for_output_layer -= self.learning_rate *
output_layer_gradient

    self.weights_from_input_to_hidden -= self.learning_rate *
input_to_hidden_weight_update
    self.bias_for_hidden_layer -= self.learning_rate *
hidden_layer_gradient

def train(self, training_data, training_labels, epochs):
    for epoch in range(epochs):
        total_loss = 0
        for inputs, actual_output in zip(training_data,
training_labels):
            predicted_output = self.forward_prop(inputs)
            self.backward_prop(inputs, actual_output,
predicted_output)
            total_loss += (actual_output - predicted_output) ** 2

        if (epoch + 1) % 100 == 0:
            print(f"Epoch {epoch + 1}/{epochs}, Loss: {total_loss
/ len(training_data)}")

def predict(self, new_data_record):
    return self.forward_prop(new_data_record)

def calc_error(self, X_test, y_test):
    # Make predictions on the test set
    y_pred = np.array([self.predict(x) for x in X_test])

    # Calculate mean squared error (MSE)
    mse = np.mean((y_test - y_pred.flatten()) ** 2)
    return mse

```

Third: Let's train our Neural Network

Now, time to train the network using the training data for a specified number of epochs.

```

# Define the neural network architecture
input_size = 4 # Number of features
hidden_size = 10 # Number of neurons in the hidden layer (can be
adjusted and tuned)
output_size = 1 # Single output (concrete strength)
learning_rate = 0.001

```

```
# Initialize our neural network
neural_network = NeuralNetwork(input_size, hidden_size, output_size,
learning_rate)

# Train the network
epochs = 1000 # Number of training epochs
neural_network.train(X_train, y_train, epochs)

Epoch 100/1000, Loss: [65.20626636]
Epoch 200/1000, Loss: [59.67481446]
Epoch 300/1000, Loss: [57.3332437]
Epoch 400/1000, Loss: [55.54652724]
Epoch 500/1000, Loss: [53.42365703]
Epoch 600/1000, Loss: [51.76129296]
Epoch 700/1000, Loss: [50.95980553]
Epoch 800/1000, Loss: [50.15264839]
Epoch 900/1000, Loss: [49.46731329]
Epoch 1000/1000, Loss: [48.97239151]
```

Finally: Let's evaluate and predict

We're going to Use the trained network to make predictions and evaluate its performance on the test set.

```
# Calculate mean squared error (MSE)
mse = neural_network.calc_error(X_test, y_test)
print(f"Mean Squared Error on Test Set: {mse}")

Mean Squared Error on Test Set: 49.064507480434074
```